

```
In [2]:
```

```
import pandas as pd
import numpy as np
from sklearn import preprocessing
import warnings
warnings.filterwarnings('ignore')
```

```
In [3]:
```

```
data = pd.read_csv("Churn.csv",na_values='NaN')
data
```

```
Out[3]:
```

	Unnamed: 0	state	area.code	account.length	voice.plan	voice.messages	intl.plan	intl.mins	intl.calls	intl.charge	...	da
0	1	KS	area_code_415	128	yes	25	no	10.0	3	2.70	...	
1	2	OH	area_code_415	107	yes	26	no	13.7	3	3.70	...	
2	3	NJ	area_code_415	137	no	0	no	12.2	5	3.29	...	
3	4	OH	area_code_408	84	no	0	yes	6.6	7	1.78	...	
4	5	OK	area_code_415	75	no	0	yes	10.1	3	2.73	...	
...	
4995	4996	HI	area_code_408	50	yes	40	no	9.9	5	2.67	...	
4996	4997	WV	area_code_415	152	no	0	no	14.7	2	3.97	...	
4997	4998	DC	area_code_415	61	no	0	no	13.6	4	3.67	...	
4998	4999	DC	area_code_510	109	no	0	no	8.5	6	2.30	...	
4999	5000	VT	area_code_415	86	yes	34	no	9.3	16	2.51	...	

5000 rows × 21 columns

```
In [4]:
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   Unnamed: 0        5000 non-null   int64  
 1   state            5000 non-null   object  
 2   area.code         5000 non-null   object  
 3   account.length   5000 non-null   int64  
 4   voice.plan       5000 non-null   object  
 5   voice.messages   5000 non-null   int64  
 6   intl.plan        5000 non-null   object  
 7   intl.mins        5000 non-null   float64 
 8   intl.calls       5000 non-null   int64  
 9   intl.charge      5000 non-null   float64 
 10  day.mins         5000 non-null   float64 
 11  day.calls        5000 non-null   int64  
 12  day.charge       4993 non-null   float64 
 13  eve.mins         4976 non-null   float64 
 14  eve.calls        5000 non-null   int64  
 15  eve.charge       5000 non-null   float64 
 16  night.mins       5000 non-null   float64 
 17  night.calls      5000 non-null   int64  
 18  night.charge     5000 non-null   float64 
 19  customer.calls   5000 non-null   int64  
 20  churn            5000 non-null   object  
dtypes: float64(8), int64(8), object(5)
memory usage: 820.4+ KB
```

```
In [5]:
```

```
data.isna().sum()
```

```
Out[5]: Unnamed: 0      0
state          0
area.code      0
account.length 0
voice.plan     0
voice.messages 0
intl.plan      0
intl.mins      0
intl.calls     0
intl.charge    0
day.mins       0
day.calls      0
day.charge     7
eve.mins       24
eve.calls      0
eve.charge     0
night.mins     0
night.calls    0
night.charge   0
customer.calls 0
churn          0
dtype: int64
```

```
In [6]: data.dropna(inplace=True)
```

```
In [7]: data.columns
```

```
Out[7]: Index(['Unnamed: 0', 'state', 'area.code', 'account.length', 'voice.plan',
               'voice.messages', 'intl.plan', 'intl.mins', 'intl.calls', 'intl.charge',
               'day.mins', 'day.calls', 'day.charge', 'eve.mins', 'eve.calls',
               'eve.charge', 'night.mins', 'night.calls', 'night.charge',
               'customer.calls', 'churn'],
              dtype='object')
```

```
In [8]: data.rename(columns={'area.code':'area_code','account.length':'account_length','voice.plan':'voice_plan',
                           'voice.messages':'voice_messages','intl.plan':'intl_plan','intl.mins':'intl_mins',
                           'intl.calls':'intl_calls','intl.charge':'intl_charge','day.mins':'day_mins','day.calls':'day_calls',
                           'day.charge':'day_charge','eve.mins':'eve_mins','eve.calls':'eve_calls','eve.charge':'eve_charge',
                           'night.mins':'night_mins','night.calls':'night_calls','night.charge':'night_charge','customer.calls':'customer_calls'},
                           inplace=True)
data
```

```
Out[8]:   Unnamed: 0  state  area_code  account_length  voice_plan  voice_messages  intl_plan  intl_mins  intl_calls  intl_charge ...
0           1    KS  area_code_415            128      yes           25        no     10.0       3      2.70 ...
1           2    OH  area_code_415            107      yes           26        no     13.7       3      3.70 ...
2           3    NJ  area_code_415            137       no            0        no     12.2       5      3.29 ...
3           4    OH  area_code_408             84       no            0        yes      6.6       7      1.78 ...
4           5    OK  area_code_415             75       no            0        yes     10.1       3      2.73 ...
...
4995        4996   HI  area_code_408            50      yes           40        no      9.9       5      2.67 ...
4996        4997   WV  area_code_415            152       no            0        no     14.7       2      3.97 ...
4997        4998   DC  area_code_415             61       no            0        no     13.6       4      3.67 ...
4998        4999   DC  area_code_510            109       no            0        no      8.5       6      2.30 ...
4999        5000   VT  area_code_415             86      yes           34        no      9.3      16      2.51 ...
```

4969 rows × 21 columns

```
In [9]: data.drop('Unnamed: 0',axis=1,inplace=True)
data
```

Out[9]:

	state	area_code	account_length	voice_plan	voice_messages	intl_plan	intl_mins	intl_calls	intl_charge	day_mins	da
0	KS	area_code_415	128	yes	25	no	10.0	3	2.70	265.1	
1	OH	area_code_415	107	yes	26	no	13.7	3	3.70	161.6	
2	NJ	area_code_415	137	no	0	no	12.2	5	3.29	243.4	
3	OH	area_code_408	84	no	0	yes	6.6	7	1.78	299.4	
4	OK	area_code_415	75	no	0	yes	10.1	3	2.73	166.7	
...
4995	HI	area_code_408	50	yes	40	no	9.9	5	2.67	235.7	
4996	WV	area_code_415	152	no	0	no	14.7	2	3.97	184.2	
4997	DC	area_code_415	61	no	0	no	13.6	4	3.67	140.6	
4998	DC	area_code_510	109	no	0	no	8.5	6	2.30	188.8	
4999	VT	area_code_415	86	yes	34	no	9.3	16	2.51	129.4	

4969 rows × 20 columns

In [10]:

```
le = preprocessing.LabelEncoder()
objlist = ['area_code','intl_plan','voice_plan','churn','state']
data[objlist] = data[objlist].apply(le.fit_transform)
```

In [11]:

```
data
```

Out[11]:

	state	area_code	account_length	voice_plan	voice_messages	intl_plan	intl_mins	intl_calls	intl_charge	day_mins	day_ca
0	16	1	128	1	25	0	10.0	3	2.70	265.1	1
1	35	1	107	1	26	0	13.7	3	3.70	161.6	1
2	31	1	137	0	0	0	12.2	5	3.29	243.4	1
3	35	0	84	0	0	1	6.6	7	1.78	299.4	
4	36	1	75	0	0	1	10.1	3	2.73	166.7	1
...
4995	11	0	50	1	40	0	9.9	5	2.67	235.7	1
4996	49	1	152	0	0	0	14.7	2	3.97	184.2	
4997	7	1	61	0	0	0	13.6	4	3.67	140.6	
4998	7	2	109	0	0	0	8.5	6	2.30	188.8	
4999	46	1	86	1	34	0	9.3	16	2.51	129.4	1

4969 rows × 20 columns

In [12]:

```
data.columns
```

Out[12]:

```
Index(['state', 'area_code', 'account_length', 'voice_plan', 'voice_messages',
       'intl_plan', 'intl_mins', 'intl_calls', 'intl_charge', 'day_mins',
       'day_calls', 'day_charge', 'eve_mins', 'eve_calls', 'eve_charge',
       'night_mins', 'night_calls', 'night_charge', 'customer_calls', 'churn'],
      dtype='object')
```

In [13]:

```
data = data[['account_length', 'voice_plan', 'voice_messages',
            'intl_plan', 'intl_calls', 'intl_charge',
            'day_calls', 'day_charge', 'eve_calls', 'eve_charge',
            'night_calls', 'night_charge', 'customer_calls', 'churn']]
```

data

Out[13]:

	account_length	voice_plan	voice_messages	intl_plan	intl_calls	intl_charge	day_calls	day_charge	eve_calls	eve_charge
0	128	1	25	0	3	2.70	110	45.07	99	16.78
1	107	1	26	0	3	3.70	123	27.47	103	16.62
2	137	0	0	0	5	3.29	114	41.38	110	10.30
3	84	0	0	1	7	1.78	71	50.90	88	5.26
4	75	0	0	1	3	2.73	113	28.34	122	12.61
...
4995	50	1	40	0	5	2.67	127	40.07	126	18.96
4996	152	0	0	0	2	3.97	90	31.31	73	21.83
4997	61	0	0	0	4	3.67	89	23.90	128	14.69
4998	109	0	0	0	6	2.30	67	32.10	92	14.59
4999	86	1	34	0	16	2.51	102	22.00	104	22.70

4969 rows × 14 columns

```
In [14]: data['Total_Charge'] = data['day_charge'] + data['eve_charge'] + data['night_charge']
```

```
In [15]: data
```

Out[15]:

	account_length	voice_plan	voice_messages	intl_plan	intl_calls	intl_charge	day_calls	day_charge	eve_calls	eve_charge
0	128	1	25	0	3	2.70	110	45.07	99	16.78
1	107	1	26	0	3	3.70	123	27.47	103	16.62
2	137	0	0	0	5	3.29	114	41.38	110	10.30
3	84	0	0	1	7	1.78	71	50.90	88	5.26
4	75	0	0	1	3	2.73	113	28.34	122	12.61
...
4995	50	1	40	0	5	2.67	127	40.07	126	18.96
4996	152	0	0	0	2	3.97	90	31.31	73	21.83
4997	61	0	0	0	4	3.67	89	23.90	128	14.69
4998	109	0	0	0	6	2.30	67	32.10	92	14.59
4999	86	1	34	0	16	2.51	102	22.00	104	22.70

4969 rows × 15 columns

```
In [16]: data.drop(['day_charge','eve_charge','night_charge'],axis=1,inplace=True)
```

```
In [17]: data=data.iloc[:,[10,0,1,2,3,4,5,6,7,8,9,11]]
```

```
In [18]: data
```

Out[18]:

	churn	account_length	voice_plan	voice_messages	intl_plan	intl_calls	intl_charge	day_calls	eve_calls	night_calls	custc
0	0	128	1	25	0	3	2.70	110	99	91	
1	0	107	1	26	0	3	3.70	123	103	103	
2	0	137	0	0	0	5	3.29	114	110	104	
3	0	84	0	0	1	7	1.78	71	88	89	
4	0	75	0	0	0	1	2.73	113	122	121	
...
4995	0	50	1	40	0	5	2.67	127	126	116	
4996	1	152	0	0	0	2	3.97	90	73	113	
4997	0	61	0	0	0	4	3.67	89	128	97	
4998	0	109	0	0	0	6	2.30	67	92	89	
4999	0	86	1	34	0	16	2.51	102	104	100	

4969 rows × 12 columns

```
In [19]: x = data.iloc[:,1:]  
x
```

```
Out[19]:   account_length  voice_plan  voice_messages  intl_plan  intl_calls  intl_charge  day_calls  eve_calls  night_calls  customer_c  
          0             128           1            25         0            3        2.70       110         99         91  
          1             107           1            26         0            3        3.70       123        103        103  
          2             137           0            0         0            5        3.29       114        110        104  
          3              84           0            0         0            1        1.78        71         88         89  
          4              75           0            0         0            1        2.73       113        122        121  
          ...           ...           ...           ...           ...           ...           ...           ...           ...           ...           ...  
          4995            50           1            40         0            5        2.67       127        126        116  
          4996            152           0            0         0            2        3.97       90         73        113  
          4997            61           0            0         0            4        3.67       89        128        97  
          4998            109           0            0         0            6        2.30       67         92        89  
          4999            86           1            34         0           16        2.51      102        104        100
```

4969 rows × 11 columns

```
In [20]: y = data.iloc[:,0:1]  
y
```

```
Out[20]:   churn  
          0     0  
          1     0  
          2     0  
          3     0  
          4     0  
          ...   ...  
          4995    0  
          4996    1  
          4997    0  
          4998    0  
          4999    0
```

4969 rows × 1 columns

```
In [21]: from sklearn.preprocessing import RobustScaler
```

```
In [22]: transformer = RobustScaler().fit(x)  
x=transformer.transform(x)
```

```
In [23]: x
```

```
Out[23]: array([[ 0.51851852,  1.          ,  1.47058824,  ..., -0.34615385,  
                  0.          ,  1.16008615],  
                 [ 0.12962963,  1.          ,  1.52941176,  ...,  0.11538462,  
                  0.          , -0.08327351],  
                 [ 0.68518519,  0.          ,  0.          ,  ...,  0.15384615,  
                  -1.          ,  0.16511127],  
                 ...,  
                 [-0.72222222,  0.          ,  0.          ,  ..., -0.11538462,  
                  0.          , -0.6137832 ],  
                 [ 0.16666667,  0.          ,  0.          ,  ..., -0.42307692,  
                  -1.          ,  0.00646088],  
                 [-0.25925926,  1.          ,  2.          ,  ...,  0.          ,  
                  -1.          , -0.36109117]])
```

```
In [24]: from sklearn.model_selection import train_test_split, cross_val_score
```

```
In [25]: x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3, random_state=15)
```

```
In [26]: #Logistic Regression  
from sklearn.linear_model import LogisticRegression  
from sklearn.metrics import classification_report  
from sklearn import preprocessing
```

```
from sklearn import metrics
```

```
In [27]: model=LogisticRegression()
model.fit(x_train,y_train)
y_pred=model.predict(x_test)
```

```
In [28]: from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score as ac
confusion_matrix=confusion_matrix(y_test,y_pred)
print(confusion_matrix)
```

```
[1254 39]
[ 156 42]
```

```
In [29]: print(classification_report(y_test,y_pred))
from sklearn.metrics import precision_recall_fscore_support as score
precision,recall,fscore,support = score(y_test,y_pred,average='macro')
LR=fscore
LR
```

	precision	recall	f1-score	support
0	0.89	0.97	0.93	1293
1	0.52	0.21	0.30	198
accuracy			0.87	1491
macro avg	0.70	0.59	0.61	1491
weighted avg	0.84	0.87	0.84	1491

```
Out[29]: 0.6144666022221427
```

```
In [30]: #KNN
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.neighbors import KNeighborsClassifier
import matplotlib.pyplot as plt
%matplotlib inline
k_range = [2*i+1 for i in range(0,20)]
k_scores = []
for k in k_range:
    knn = KNeighborsClassifier(n_neighbors=k)
    scores = cross_val_score(knn,x_train , y_train, cv = 10)
    k_scores.append(scores.mean())
    print("K value=",k)
    model = KNeighborsClassifier(n_neighbors=k)
    model.fit(x_train,y_train)
    pred = model.predict(x_test)
    print(classification_report(y_test,pred))
```

K value= 1

	precision	recall	f1-score	support
0	0.92	0.94	0.93	1293
1	0.56	0.48	0.52	198
accuracy			0.88	1491
macro avg	0.74	0.71	0.73	1491
weighted avg	0.87	0.88	0.88	1491

K value= 3

	precision	recall	f1-score	support
0	0.91	0.98	0.95	1293
1	0.75	0.40	0.52	198
accuracy			0.90	1491
macro avg	0.83	0.69	0.73	1491
weighted avg	0.89	0.90	0.89	1491

K value= 5

	precision	recall	f1-score	support
0	0.91	0.99	0.95	1293
1	0.84	0.38	0.52	198
accuracy			0.91	1491
macro avg	0.88	0.68	0.74	1491
weighted avg	0.90	0.91	0.89	1491

K value= 7

	precision	recall	f1-score	support
0	0.91	0.99	0.95	1293

1	0.90	0.33	0.49	198
accuracy			0.91	1491
macro avg	0.91	0.66	0.72	1491
weighted avg	0.91	0.91	0.89	1491

K value= 9				
	precision	recall	f1-score	support
0	0.90	0.99	0.94	1293

1	0.87	0.27	0.42	198
accuracy			0.90	1491
macro avg	0.89	0.63	0.68	1491
weighted avg	0.90	0.90	0.87	1491

K value= 11				
	precision	recall	f1-score	support
0	0.90	1.00	0.94	1293

1	0.91	0.26	0.40	198
accuracy			0.90	1491
macro avg	0.90	0.63	0.67	1491
weighted avg	0.90	0.90	0.87	1491

K value= 13				
	precision	recall	f1-score	support
0	0.89	1.00	0.94	1293

1	0.88	0.22	0.35	198
accuracy			0.89	1491
macro avg	0.89	0.61	0.65	1491
weighted avg	0.89	0.89	0.86	1491

K value= 15				
	precision	recall	f1-score	support
0	0.89	1.00	0.94	1293

1	0.87	0.21	0.33	198
accuracy			0.89	1491
macro avg	0.88	0.60	0.64	1491
weighted avg	0.89	0.89	0.86	1491

K value= 17				
	precision	recall	f1-score	support
0	0.89	1.00	0.94	1293

1	0.89	0.21	0.34	198
accuracy			0.89	1491
macro avg	0.89	0.60	0.64	1491
weighted avg	0.89	0.89	0.86	1491

K value= 19				
	precision	recall	f1-score	support
0	0.89	1.00	0.94	1293

1	0.85	0.18	0.29	198
accuracy			0.89	1491
macro avg	0.87	0.59	0.62	1491
weighted avg	0.88	0.89	0.85	1491

K value= 21				
	precision	recall	f1-score	support
0	0.89	0.99	0.94	1293

1	0.83	0.17	0.28	198
accuracy			0.89	1491
macro avg	0.86	0.58	0.61	1491
weighted avg	0.88	0.89	0.85	1491

K value= 23				
	precision	recall	f1-score	support
0	0.89	0.99	0.94	1293

1	0.84	0.18	0.30	198
accuracy			0.89	1491

macro avg	0.86	0.59	0.62	1491
weighted avg	0.88	0.89	0.85	1491

K value= 25

	precision	recall	f1-score	support
0	0.89	1.00	0.94	1293
1	0.86	0.18	0.30	198
accuracy			0.89	1491
macro avg	0.87	0.59	0.62	1491
weighted avg	0.88	0.89	0.85	1491

K value= 27

	precision	recall	f1-score	support
0	0.89	1.00	0.94	1293
1	0.85	0.17	0.28	198
accuracy			0.89	1491
macro avg	0.87	0.58	0.61	1491
weighted avg	0.88	0.89	0.85	1491

K value= 29

	precision	recall	f1-score	support
0	0.88	1.00	0.94	1293
1	0.86	0.15	0.26	198
accuracy			0.88	1491
macro avg	0.87	0.57	0.60	1491
weighted avg	0.88	0.88	0.85	1491

K value= 31

	precision	recall	f1-score	support
0	0.89	1.00	0.94	1293
1	0.89	0.16	0.27	198
accuracy			0.89	1491
macro avg	0.89	0.58	0.60	1491
weighted avg	0.89	0.89	0.85	1491

K value= 33

	precision	recall	f1-score	support
0	0.89	1.00	0.94	1293
1	0.89	0.16	0.27	198
accuracy			0.89	1491
macro avg	0.89	0.58	0.60	1491
weighted avg	0.89	0.89	0.85	1491

K value= 35

	precision	recall	f1-score	support
0	0.88	1.00	0.94	1293
1	0.84	0.14	0.23	198
accuracy			0.88	1491
macro avg	0.86	0.57	0.59	1491
weighted avg	0.88	0.88	0.84	1491

K value= 37

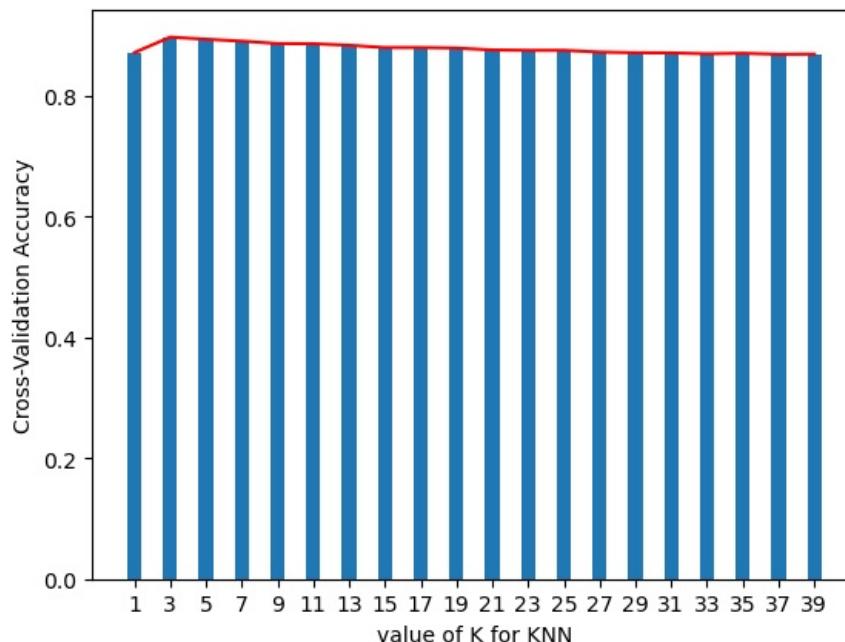
	precision	recall	f1-score	support
0	0.88	1.00	0.94	1293
1	0.87	0.14	0.24	198
accuracy			0.88	1491
macro avg	0.88	0.57	0.59	1491
weighted avg	0.88	0.88	0.84	1491

K value= 39

	precision	recall	f1-score	support
0	0.88	1.00	0.94	1293
1	0.90	0.13	0.23	198
accuracy			0.88	1491
macro avg	0.89	0.56	0.58	1491
weighted avg	0.88	0.88	0.84	1491

```
In [31]: plt.bar(k_range, k_scores)
plt.plot(k_range, k_scores, color = "red")

plt.xlabel('value of K for KNN')
plt.ylabel('Cross-Validation Accuracy')
plt.xticks(k_range)
plt.show()
```



```
In [32]: model = KNeighborsClassifier(n_neighbors=3)
model.fit(x_train,y_train)
pred = model.predict(x_test)
```

```
In [33]: print(classification_report(y_test,pred))
precision,recall,fscore,support = score(y_test,pred,average='macro')
KNN=fscore
KNN
```

	precision	recall	f1-score	support
0	0.91	0.98	0.95	1293
1	0.75	0.40	0.52	198
accuracy			0.90	1491
macro avg	0.83	0.69	0.73	1491
weighted avg	0.89	0.90	0.89	1491

```
Out[33]: 0.7336637359144649
```

```
In [34]: #Decision Tree
from sklearn.tree import DecisionTreeClassifier
from sklearn import tree
model = DecisionTreeClassifier(criterion='gini',min_samples_split=5)
model.fit(x_train,y_train)
pred = model.predict(x_test)
metrics.accuracy_score(pred,y_test)
print(classification_report(y_test,pred))
precision,recall,fscore,support = score(y_test,pred,average='macro')
DT=fscore
DT
```

	precision	recall	f1-score	support
0	0.97	0.98	0.97	1293
1	0.85	0.80	0.83	198
accuracy			0.96	1491
macro avg	0.91	0.89	0.90	1491
weighted avg	0.95	0.96	0.95	1491

```
Out[34]: 0.9000875135646025
```

```
In [35]: #Decision Tree
from sklearn.tree import DecisionTreeClassifier
from sklearn import tree
model = DecisionTreeClassifier(criterion='entropy',min_samples_split=5)
model.fit(x_train,y_train)
```

```
pred = model.predict(x_test)
metrics.accuracy_score(pred,y_test)
print(classification_report(y_test,pred))
```

	precision	recall	f1-score	support
0	0.97	0.97	0.97	1293
1	0.83	0.80	0.82	198
accuracy			0.95	1491
macro avg	0.90	0.89	0.89	1491
weighted avg	0.95	0.95	0.95	1491

In [36]: # Random Forest Classification

```
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.ensemble import RandomForestClassifier

x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25,random_state=7)

num_trees = 140
max_features = 6

model = RandomForestClassifier(n_estimators=num_trees,max_samples=0.8, max_features=max_features, random_state = 7)

model.fit(x_train,y_train)

pred = model.predict(x_test)

print(classification_report(y_test,pred))
precision,recall,fscore,support = score(y_test,pred,average='macro')
RD=fscore
RD
```

	precision	recall	f1-score	support
0	0.98	1.00	0.99	1074
1	0.98	0.85	0.91	169
accuracy			0.98	1243
macro avg	0.98	0.92	0.95	1243
weighted avg	0.98	0.98	0.98	1243

Out[36]: 0.9492445896284197

In [37]: #SVM

```
from sklearn import svm
from sklearn.svm import SVC
from sklearn.model_selection import GridSearchCV

# clf = SVC()
# param_grid = [{"kernel":["rbf"],'gamma':[50,5,10,0.5,1,0.001,0.0001,0.00001],'C':[1,15,14,13,12,11,10,0.1] }]
# gsv = GridSearchCV(clf,param_grid,cv=10)
# gsv.fit(x_train,y_train)
```

In [38]: #gsv.best_params_ , gsv.best_score_

In [39]:

```
clf = SVC(C= 1, gamma = 0.5,kernel="rbf")
clf.fit(x_train,y_train)
y_pred = clf.predict(x_test)
print(classification_report(y_test, y_pred))
precision,recall,fscore,support = score(y_test,y_pred,average='macro')
SVM=fscore
SVM
```

	precision	recall	f1-score	support
0	0.93	0.99	0.96	1074
1	0.92	0.51	0.66	169
accuracy			0.93	1243
macro avg	0.92	0.75	0.81	1243
weighted avg	0.93	0.93	0.92	1243

Out[39]: 0.8092934293429344

In [40]: #ANN

```
import tensorflow as tf
model = tf.keras.models.Sequential()
model.add(tf.keras.layers.Dense(15,input_dim=11,activation = 'sigmoid'))
model.add(tf.keras.layers.Dense(10,activation='sigmoid'))
```

```
model.add(tf.keras.layers.Dense(1, activation='sigmoid'))
```

In [41]: `model.summary()`

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 15)	180
dense_1 (Dense)	(None, 10)	160
dense_2 (Dense)	(None, 1)	11

Total params: 351 (1.37 KB)

Trainable params: 351 (1.37 KB)

Non-trainable params: 0 (0.00 B)

In [42]: `# Compile model`

```
model.compile(loss = 'binary_crossentropy', optimizer='adam', metrics=['accuracy'])
```

In [43]: `history = model.fit(x,y,validation_split=0.20, epochs=50, batch_size=100)`

```
Epoch 1/50
40/40 ━━━━━━━━━━ 2s 11ms/step - accuracy: 0.8581 - loss: 0.4515 - val_accuracy: 0.8581 - val_loss: 0.4249
Epoch 2/50
40/40 ━━━━━━━━━━ 0s 6ms/step - accuracy: 0.8581 - loss: 0.4150 - val_accuracy: 0.8581 - val_loss: 0.4078
Epoch 3/50
40/40 ━━━━━━━━━━ 0s 6ms/step - accuracy: 0.8581 - loss: 0.4055 - val_accuracy: 0.8581 - val_loss: 0.4026
Epoch 4/50
40/40 ━━━━━━━━━━ 0s 5ms/step - accuracy: 0.8581 - loss: 0.4015 - val_accuracy: 0.8581 - val_loss: 0.3990
Epoch 5/50
40/40 ━━━━━━━━━━ 0s 5ms/step - accuracy: 0.8581 - loss: 0.3983 - val_accuracy: 0.8581 - val_loss: 0.3954
Epoch 6/50
40/40 ━━━━━━━━━━ 0s 5ms/step - accuracy: 0.8581 - loss: 0.3947 - val_accuracy: 0.8581 - val_loss: 0.3917
Epoch 7/50
40/40 ━━━━━━━━━━ 0s 6ms/step - accuracy: 0.8581 - loss: 0.3910 - val_accuracy: 0.8581 - val_loss: 0.3878
Epoch 8/50
40/40 ━━━━━━━━━━ 0s 6ms/step - accuracy: 0.8581 - loss: 0.3869 - val_accuracy: 0.8581 - val_loss: 0.3833
Epoch 9/50
40/40 ━━━━━━━━━━ 0s 6ms/step - accuracy: 0.8581 - loss: 0.3826 - val_accuracy: 0.8581 - val_loss: 0.3788
Epoch 10/50
40/40 ━━━━━━━━━━ 0s 6ms/step - accuracy: 0.8581 - loss: 0.3781 - val_accuracy: 0.8581 - val_loss: 0.3741
Epoch 11/50
40/40 ━━━━━━━━━━ 0s 6ms/step - accuracy: 0.8581 - loss: 0.3734 - val_accuracy: 0.8581 - val_loss: 0.3690
Epoch 12/50
40/40 ━━━━━━━━━━ 0s 6ms/step - accuracy: 0.8581 - loss: 0.3684 - val_accuracy: 0.8581 - val_loss: 0.3637
Epoch 13/50
40/40 ━━━━━━━━━━ 0s 6ms/step - accuracy: 0.8581 - loss: 0.3633 - val_accuracy: 0.8581 - val_loss: 0.3584
Epoch 14/50
40/40 ━━━━━━━━━━ 0s 5ms/step - accuracy: 0.8581 - loss: 0.3583 - val_accuracy: 0.8581 - val_loss: 0.3532
Epoch 15/50
40/40 ━━━━━━━━━━ 0s 5ms/step - accuracy: 0.8581 - loss: 0.3534 - val_accuracy: 0.8581 - val_loss: 0.3485
Epoch 16/50
40/40 ━━━━━━━━━━ 0s 6ms/step - accuracy: 0.8581 - loss: 0.3490 - val_accuracy: 0.8581 - val_loss: 0.3437
Epoch 17/50
40/40 ━━━━━━━━━━ 0s 6ms/step - accuracy: 0.8581 - loss: 0.3446 - val_accuracy: 0.8581 - val_loss: 0.3393
Epoch 18/50
40/40 ━━━━━━━━━━ 0s 6ms/step - accuracy: 0.8581 - loss: 0.3410 - val_accuracy: 0.8581 - val_loss: 0.3354
Epoch 19/50
40/40 ━━━━━━━━━━ 0s 5ms/step - accuracy: 0.8584 - loss: 0.3376 - val_accuracy: 0.8622 - val_loss: 0.3320
Epoch 20/50
40/40 ━━━━━━━━━━ 0s 6ms/step - accuracy: 0.8594 - loss: 0.3348 - val_accuracy: 0.8632 - val_loss: 0.32
```

87
Epoch 21/50
40/40 0s 5ms/step - accuracy: 0.8601 - loss: 0.3319 - val_accuracy: 0.8652 - val_loss: 0.32
60
Epoch 22/50
40/40 0s 5ms/step - accuracy: 0.8604 - loss: 0.3296 - val_accuracy: 0.8652 - val_loss: 0.32
33
Epoch 23/50
40/40 0s 5ms/step - accuracy: 0.8634 - loss: 0.3276 - val_accuracy: 0.8612 - val_loss: 0.32
13
Epoch 24/50
40/40 0s 5ms/step - accuracy: 0.8624 - loss: 0.3262 - val_accuracy: 0.8632 - val_loss: 0.31
93
Epoch 25/50
40/40 0s 6ms/step - accuracy: 0.8626 - loss: 0.3245 - val_accuracy: 0.8632 - val_loss: 0.31
76
Epoch 26/50
40/40 0s 5ms/step - accuracy: 0.8626 - loss: 0.3234 - val_accuracy: 0.8632 - val_loss: 0.31
59
Epoch 27/50
40/40 0s 4ms/step - accuracy: 0.8636 - loss: 0.3223 - val_accuracy: 0.8622 - val_loss: 0.31
50
Epoch 28/50
40/40 0s 5ms/step - accuracy: 0.8621 - loss: 0.3215 - val_accuracy: 0.8612 - val_loss: 0.31
38
Epoch 29/50
40/40 0s 5ms/step - accuracy: 0.8636 - loss: 0.3208 - val_accuracy: 0.8632 - val_loss: 0.31
26
Epoch 30/50
40/40 0s 5ms/step - accuracy: 0.8629 - loss: 0.3202 - val_accuracy: 0.8652 - val_loss: 0.31
16
Epoch 31/50
40/40 0s 6ms/step - accuracy: 0.8667 - loss: 0.3196 - val_accuracy: 0.8652 - val_loss: 0.31
04
Epoch 32/50
40/40 0s 5ms/step - accuracy: 0.8654 - loss: 0.3191 - val_accuracy: 0.8672 - val_loss: 0.31
02
Epoch 33/50
40/40 0s 6ms/step - accuracy: 0.8674 - loss: 0.3187 - val_accuracy: 0.8662 - val_loss: 0.30
91
Epoch 34/50
40/40 0s 6ms/step - accuracy: 0.8674 - loss: 0.3181 - val_accuracy: 0.8662 - val_loss: 0.30
86
Epoch 35/50
40/40 0s 6ms/step - accuracy: 0.8657 - loss: 0.3178 - val_accuracy: 0.8662 - val_loss: 0.30
80
Epoch 36/50
40/40 0s 5ms/step - accuracy: 0.8664 - loss: 0.3175 - val_accuracy: 0.8672 - val_loss: 0.30
74
Epoch 37/50
40/40 0s 6ms/step - accuracy: 0.8667 - loss: 0.3173 - val_accuracy: 0.8652 - val_loss: 0.30
68
Epoch 38/50
40/40 0s 6ms/step - accuracy: 0.8674 - loss: 0.3172 - val_accuracy: 0.8662 - val_loss: 0.30
62
Epoch 39/50
40/40 0s 5ms/step - accuracy: 0.8657 - loss: 0.3174 - val_accuracy: 0.8632 - val_loss: 0.30
59
Epoch 40/50
40/40 0s 4ms/step - accuracy: 0.8669 - loss: 0.3166 - val_accuracy: 0.8642 - val_loss: 0.30
55
Epoch 41/50
40/40 0s 5ms/step - accuracy: 0.8669 - loss: 0.3165 - val_accuracy: 0.8632 - val_loss: 0.30
52
Epoch 42/50
40/40 0s 4ms/step - accuracy: 0.8672 - loss: 0.3163 - val_accuracy: 0.8632 - val_loss: 0.30
46
Epoch 43/50
40/40 0s 4ms/step - accuracy: 0.8662 - loss: 0.3161 - val_accuracy: 0.8632 - val_loss: 0.30
44
Epoch 44/50
40/40 0s 4ms/step - accuracy: 0.8667 - loss: 0.3160 - val_accuracy: 0.8632 - val_loss: 0.30
42
Epoch 45/50
40/40 0s 3ms/step - accuracy: 0.8664 - loss: 0.3158 - val_accuracy: 0.8642 - val_loss: 0.30
36
Epoch 46/50
40/40 0s 3ms/step - accuracy: 0.8667 - loss: 0.3155 - val_accuracy: 0.8632 - val_loss: 0.30
33
Epoch 47/50
40/40 0s 5ms/step - accuracy: 0.8659 - loss: 0.3154 - val_accuracy: 0.8632 - val_loss: 0.30
30
Epoch 48/50

```
40/40 ━━━━━━━━━━ 0s 5ms/step - accuracy: 0.8662 - loss: 0.3154 - val_accuracy: 0.8612 - val_loss: 0.30  
29  
Epoch 49/50  
40/40 ━━━━━━━━━━ 0s 5ms/step - accuracy: 0.8669 - loss: 0.3149 - val_accuracy: 0.8632 - val_loss: 0.30  
23  
Epoch 50/50  
40/40 ━━━━━━━━━━ 0s 5ms/step - accuracy: 0.8662 - loss: 0.3149 - val_accuracy: 0.8622 - val_loss: 0.30  
21
```

```
In [44]: model.save_weights("mywt.weights.h5")
```

```
In [45]: #evaluate the model  
scores = model.evaluate(x,y)  
print("%s: %.2f%%" % (model.metrics_names[1], scores[1]*100))  
ANN=model.metrics_names[1], scores[1]*100  
ANN=ANN[1]  
ANN
```

```
156/156 ━━━━━━━━━━ 0s 2ms/step - accuracy: 0.8658 - loss: 0.3120  
compile_metrics: 86.58%
```

```
Out[45]: 86.5767776966095
```

```
In [46]: #Naive_Bayes  
from sklearn.naive_bayes import GaussianNB  
nb = GaussianNB()  
nb.fit(x_train, y_train)
```

```
Out[46]: ▾ GaussianNB ⓘ ?  
GaussianNB()
```

```
In [47]: y_pred = nb.predict(x_test)
```

```
In [48]: from sklearn.metrics import confusion_matrix  
conf_matrix1=confusion_matrix(y_test, y_pred)  
conf_matrix1
```

```
Out[48]: array([[980,  94],  
                 [ 91,  78]])
```

```
In [49]: import seaborn as sns  
%matplotlib inline  
fig, ax=plt.subplots(1,1, figsize=(4,4), dpi=100)  
  
xlabels=['0-Less than 50K', '1-Above 50K']  
ylabels=['0-Less than 50K', '1-Above 50K']  
  
sns.heatmap(conf_matrix1, annot=True, cbar=None, cmap="Purples", xticklabels=xlabels, yticklabels=ylabels,  
            linewidths=1, linecolor='black')  
ax.set_xlabel('Predicted Class')  
ax.set_ylabel('True Class')
```

```
Out[49]: Text(20.72222222222214, 0.5, 'True Class')
```

```
In [50]: from sklearn.metrics import classification_report  
names=['0-Less than 50K', '1-Above 50K']  
print(classification_report(y_test, y_pred, target_names=names))  
precision, recall, fscore, support = score(y_test, y_pred, average='macro')  
NB=fscore  
NB
```

	precision	recall	f1-score	support
0-Less than 50K	0.92	0.91	0.91	1074
1-Above 50K	0.45	0.46	0.46	169
accuracy			0.85	1243
macro avg	0.68	0.69	0.69	1243
weighted avg	0.85	0.85	0.85	1243

```
Out[50]: 0.6856154598090082
```

```
In [51]: #Bagging
```

```
from sklearn.model_selection import KFold  
from sklearn.model_selection import cross_val_score  
from sklearn.ensemble import BaggingClassifier  
from sklearn.model_selection import train_test_split  
from sklearn.metrics import classification_report  
  
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=2)
```

```

num_trees = 100
model = BaggingClassifier(max_samples=0.8, n_estimators=num_trees, random_state=8)

model.fit(x_train,y_train)
pred = model.predict(x_test)
print(classification_report(y_test,pred))
precision,recall,fscore,support = score(y_test,pred,average='macro')
BG=fscore
BG

```

	precision	recall	f1-score	support
0	0.98	0.99	0.99	1290
1	0.95	0.88	0.91	201
accuracy			0.98	1491
macro avg	0.97	0.93	0.95	1491
weighted avg	0.98	0.98	0.98	1491

Out[51]: 0.9494100130132448

```

In [52]: from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import BaggingClassifier
from sklearn.model_selection import GridSearchCV

param_grid = {
    'estimator__max_depth': [1, 2, 3, 4, 5], # FIXED HERE
    'max_samples': [0.05, 0.1, 0.2, 0.5]
}

bag = BaggingClassifier(
    estimator=DecisionTreeClassifier(), # FIXED HERE
    n_estimators=100,
    max_features=0.5,
    random_state=42
)

clf = GridSearchCV(
    bag,
    param_grid,
    cv=5,
    n_jobs=-1
)

clf.fit(x_train, y_train)

```

Out[52]:

```

        ► GridSearchCV
            ⓘ ⓘ
        ► best_estimator_: BaggingClassifier
            ► estimator:
                DecisionTreeClassifier
                    ► DecisionTreeClassifier ⓘ

```

In [53]: clf.best_params_, clf.best_score_

Out[53]: ({'estimator__max_depth': 5, 'max_samples': 0.5},
np.float64(0.8723397833457373))

In [54]: #AdaBoost Classifier

```

from sklearn.ensemble import AdaBoostClassifier
from sklearn.tree import DecisionTreeClassifier

ada = AdaBoostClassifier(
    estimator=DecisionTreeClassifier(max_depth=1), # weak learner
    n_estimators=200,
    learning_rate=0.05,
    random_state=42
)

ada.fit(x_train, y_train)

y_pred = ada.predict(x_test)
from sklearn.metrics import accuracy_score

#Accuracy (quick check)
acc = accuracy_score(y_test, y_pred)

```

```

print("Accuracy:", acc)

#Confusion Matrix (MOST IMPORTANT)
from sklearn.metrics import confusion_matrix

cm = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:\n", cm)

import seaborn as sns
import matplotlib.pyplot as plt

sns.heatmap(cm, annot=True, fmt="d", cmap="Blues")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()

#Precision, Recall, F1-score (INTERVIEW FAVORITE)
from sklearn.metrics import classification_report

print(classification_report(y_test, y_pred))

#ROC-AUC Score (BEST METRIC for churn)
from sklearn.metrics import roc_auc_score

y_prob = ada.predict_proba(x_test)[:, 1]
roc_auc = roc_auc_score(y_test, y_prob)

print("ROC-AUC Score:", roc_auc)

```

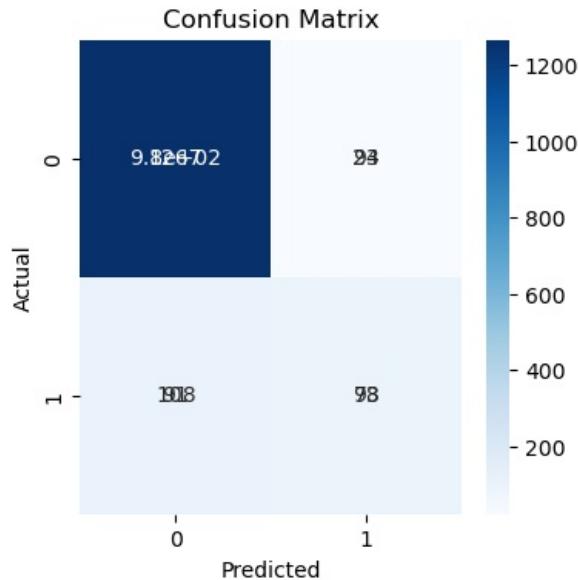
Accuracy: 0.9121395036887995

Confusion Matrix:

```

[[1267  23]
 [ 108  93]]

```



	precision	recall	f1-score	support
0	0.92	0.98	0.95	1290
1	0.80	0.46	0.59	201
accuracy			0.91	1491
macro avg	0.86	0.72	0.77	1491
weighted avg	0.91	0.91	0.90	1491

ROC-AUC Score: 0.9133036368544872

Accuracy Scores All Models

```

In [55]: # initialize list elements
data1 = [['Logistic Regression',LR],['KNearest Nighbour',KNN],
         ['Decision Tree',DT],['Random Forest',RD],['Support Vector Machine',SVM],['Artificial Neural Network',ANN],
         ['Navie Bais',NB],['Bagging',BG],['AdaBoosting',ada]]
# Create the pandas DataFrame with column name is provided explicitly
df = pd.DataFrame(data1, columns=['Algorithm Names','Accuracy'])
df

```

Out[55]:

	Algorithm Names	Accuracy
0	Logistic Regression	0.614467
1	KNearest Nighbour	0.733664
2	Decision Tree	0.900088
3	Random Forest	0.949245
4	Support Vector Machine	0.809293
5	Artificial Neural Network	86.576778
6	Navie Bais	0.685615
7	Bagging	0.94941
8	AdaBoosting (DecisionTreeClassifier(max_depth=1, random_st...	

In [56]:

```
y_test = y_test.to_numpy()
y_test
```

Out[56]:

```
array([[0],
       [0],
       [0],
       ...,
       [0],
       [0],
       [0]])
```

In [60]:

```
import numpy as np
from sklearn.neural_network import MLPClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score

x_train, x_test, y_train, y_test = train_test_split(
    x, y, test_size=0.3, random_state=15
)

clf = MLPClassifier(
    hidden_layer_sizes=(10,),
    max_iter=1000,
    random_state=15
)

clf.fit(x_train, y_train)

y_pred = clf.predict(x_test)

print(len(y_pred))
print(len(y_test))
print(y_pred)

accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy * 100:.2f}%")
```

```
1491
1491
[0 0 1 ... 0 0 0]
Accuracy: 94.43%
```

In [58]:

```
data
```

Out[58]:

	churn	account_length	voice_plan	voice_messages	intl_plan	intl_calls	intl_charge	day_calls	eve_calls	night_calls	custc
0	0	128	1	25	0	3	2.70	110	99	91	
1	0	107	1	26	0	3	3.70	123	103	103	
2	0	137	0	0	0	5	3.29	114	110	104	
3	0	84	0	0	1	7	1.78	71	88	89	
4	0	75	0	0	1	3	2.73	113	122	121	
...
4995	0	50	1	40	0	5	2.67	127	126	116	
4996	1	152	0	0	0	2	3.97	90	73	113	
4997	0	61	0	0	0	4	3.67	89	128	97	
4998	0	109	0	0	0	6	2.30	67	92	89	
4999	0	86	1	34	0	16	2.51	102	104	100	

4969 rows × 12 columns

In [59]:

x

Out[59]:

```
array([[ 0.51851852,  1.          ,  1.47058824,  ..., -0.34615385,
        0.          ,  1.16008615],
       [ 0.12962963,  1.          ,  1.52941176,  ...,  0.11538462,
        0.          , -0.08327351],
       [ 0.68518519,  0.          ,  0.          ,  ...,  0.15384615,
        -1.         ,  0.16511127],
       ...,
      [-0.72222222,  0.          ,  0.          ,  ..., -0.11538462,
        0.          , -0.6137832 ],
       [ 0.16666667,  0.          ,  0.          ,  ..., -0.42307692,
        -1.         ,  0.00646088],
      [-0.25925926,  1.          ,  2.          ,  ...,  0.          ,
        -1.         , -0.36109117]])
```

In [61]:

y

Out[61]:

	churn
0	0
1	0
2	0
3	0
4	0
...	...
4995	0
4996	1
4997	0
4998	0
4999	0

	churn
0	0
1	0
2	0
3	0
4	0
...	...
4995	0
4996	1
4997	0
4998	0
4999	0

	churn
0	0
1	0
2	0
3	0
4	0
...	...
4995	0
4996	1
4997	0
4998	0
4999	0

	churn
0	0
1	0
2	0
3	0
4	0
...	...
4995	0
4996	1
4997	0
4998	0
4999	0

	churn
0	0
1	0
2	0
3	0
4	0
...	...
4995	0
4996	1
4997	0
4998	0
4999	0

	churn
0	0
1	0
2	0
3	0
4	0
...	...
4995	0
4996	1
4997	0
4998	0
4999	0

	churn
0	0
1	0
2	0
3	0
4	0
...	...
4995	0
4996	1
4997	0
4998	0
4999	0

	churn
0	0
1	0
2	0
3	0
4	0
...	...
4995	0
4996	1
4997	0
4998	0
4999	0

	churn
0	0
1	0
2	0
3	0
4	0
...	...
4995	0
4996	1
4997	0
4998	0
4999	0

	churn
0	0
1	0
2	0
3	0
4	0
...	...
4995	0
4996	1
4997	0
4998	0
4999	0

	churn
0	0
1	0
2	0
3	0
4	0
...	...
4995	0
4996	1
4997	0
4998	0
4999	0

	churn
0	0
1	0
2	0
3	0
4	0
...	...
4995	0
4996	1
4997	0
4998	0
4999	0

	churn
0	0
1	0
2	0
3	0
4	0
...	...
4995	0
4996	1
4997	0
4998	0
4999	0

	churn
0	0
1	0
2	0
3	0
4	0
...	...
4995	0
4996	1
4997	0
4998	0
4999	0

	churn
0	0
1	0
2	0
3	0
4	0
...	...
4995	0
4996	1
4997	0
4998	0
4999	0

	churn
0	0
1	0
2	0
3	0
4	0
...	...
4995	0
4996	1
4997	0
4998	0
4999	0

	churn
0	0
1	0
2	0
3	0
4	0
...	...
4995	0
4996	1
4997	0
4998	0
4999	0

	churn
0	0
1	0
2	0
3	0
4	0
...	...
4995	0
4996	1
4997	0
4998	0
4999	0

	churn
0	0
1	0
2	0
3	0
4	0
...	...
4995	0
4996	1
4997	0
4998	0
4999	0

	churn
0	0
1	0
2	0
3	0
4	0
...	...
4995	0
4996	1
4997	0
4998	0
4999	0

	churn
0	0
1	0
2	0
3	0
4	0
...	...
4995	0
4996	1
4997	0
4998	0
4999	0

```
# Calculate the accuracy of the classifier
accuracy = np.mean(y_pred == y_test)
print("Accuracy: {:.2f}%".format(accuracy * 100))
```

In [63]:

```
import pickle

final_artifacts = {
    "model": model,                      # trained Random Forest
    "scaler": transformer,                # RobustScaler
    "label_encoder": le,                  # LabelEncoder
    "feature_order": [
        'account_length',
        'voice_plan',
        'voice_messages',
        'intl_plan',
        'intl_calls',
        'intl_charge',
        'day_calls',
        'eve_calls',
        'night_calls',
        'customer_calls',
        'Total_Charge'
    ]
}

with open("churn_complete_pipeline.pkl", "wb") as f:
    pickle.dump(final_artifacts, f)

print("✓ Full project pipeline saved successfully!")
```

✓ Full project pipeline saved successfully!

In [65]:

```
import os
print(os.getcwd())
print(os.listdir())
```

```
C:\Users\Agnel Sharon Jerald\Machine learning
['.ipynb_checkpoints', 'boosting.ipynb', 'Cars.csv', 'Churn.csv', 'churn_complete_pipeline.pkl', 'churn_coustome_r_prediction.zip', 'CHURN_CUSTOMER_PREDICTION_WITH FEATURE_ENGINEERING.ipynb', 'CHURN_CUSTOMER_WITHOUT FEATURE_E NGINEERING.ipynb', 'claimants.csv', 'Data Visualisations.ipynb', 'Decision Tree Classifier.ipynb', 'Decision Tre e_C5.0 CART.ipynb', 'Iris Species decicision Tree.ipynb', 'KNN_updated.ipynb', 'Linear_Regression.ipynb', 'Logis tic Regression.ipynb', 'Machine Learning basics.ipynb', 'Multi Linear Regression.ipynb', 'mywt.weights.h5', 'Num py.ipynb', 'pima-indians-diabetes.csv', 'polynomial Regression.ipynb', 'Practice libraries.ipynb', 'Predicting C ustomer Purchase Behavior in the Travel Industry.ipynb', 'Random Forest Regression Implementation.ipynb', 'Trave l.csv']
```

In []: